

# Secure Programming Lab 1: Data Corruption

School of Informatics, University of Edinburgh

2pm-5.30pm, 31st January 2014

This is the first laboratory session out of four in the Secure Programming course. Each lab session will examine some software vulnerabilities, including their exploits and repairs. The lab is split into several exercises.

Lab exercise are part of the delivery of the Secure Programming course. Your work in labs does not contribute towards the final mark for the course, but you are expected to take part. You will need to do this to gain the background needed for the upcoming coursework assignment that *does* contribute towards the final mark, as well develop knowledge that will be required in the exam.

- **Guided introduction.** The lab will start with a short introduction, giving some hints about the exercises. Please arrive on time for this.
- **Tools and techniques.** There are some pointers to useful tools in the exercises and in the appendix at the end of this handout.
- **Working together is encouraged.** We want to foster a supportive learning environment. Students who have prior knowledge or expertise are especially encouraged to work with others. Collaborating on the exercises may help you to think more deeply about the problems by discussing different aspects, as well as sharing existing knowledge.
- **Course staff will be on hand.** We will be here to discuss your progress and help with solving the problems. Detailed help will only be available during the timetabled labs.
- **Submit answers.** There are **checkpoint questions** in each exercise which you can answer to measure your progress. You may submit answers using the electronic submission mechanism or simply discuss them in labs. Submitting answers will allow us to give feedback at the next session.

As the work is intended to be completed in the lab and not take more of your time, there will be a short deadline for submissions. For this lab session, submissions will be accepted until **4pm, Monday 3rd February**.

For an online copy of this handout, visit  
<http://www.inf.ed.ac.uk/teaching/courses/sp/2013/labs/lab1/>.

## Virtual machinery

For this lab we are providing some code and a virtual machine for you to use. The VM has two users, **user** and **root**. The passwords are the same as their usernames. To install the VM, you should use a virtual disk file stored in the scratch space on your machine, and let Virtual Box use it from there.

To do this you will need to change some settings in VirtualBox. First set Settings→General→Default Machine Folder to

```
/disk/scratch/module-sp/[your username]
```

Then import the appliance from the file:

```
/afs/inf.ed.ac.uk/group/teaching/module-sp/SecureProgramming-1.ova
```

The machine is set to use NAT. Once started you can either use the console window, or SSH in via your local machine over port 2222 (recommended).

```
ssh -p 2222 user@localhost
```

Lab exercises are in the folders in `/home/user/` on the VM. To compile the examples, just type **make**.

Beware that `/disk/scratch` is local to your current workstation, it is not backed up. So you should save any work that you do inside the virtual machine (edited source files, etc) in your home directory. We also recommend keeping a **lab notebook** with notes of what you have done during the exercise.

We've supplied some tools to make things easier but feel free to install additional software in the VM. The package manager is called **yaourt**.

## Driving without brakes

The machines (and exercises) have been configured to disable some preventative security measures. This makes exploits easier, so you can try out some of the classical exploitation techniques. For example, ASLR is disabled in the Linux kernel of our VM (`/proc/sys/kernel/randomize_va_space = 0`).

On a modern Linux system classic overflow attacks should be hard or impossible. But some cut-down operating systems and embedded systems are either still vulnerable or have only recently implemented protection. In Android, for example, before version 4 [ASLR wasn't used](#) and it [wasn't fully enabled until 4.1](#). As another example, [Show 93 of Gary McGraw's Silver Bullet Security Podcast](#) describes remotely exploiting a car's brake system by using multiple buffer overflows!

## Exercise 1

The VM contains a simple *Notice Board* system for a shared machine. It runs as the root user and allows any other user on the machine to append notes to a notice board file which has permissions `-rw-r--r--`. The program is written in C. See the Appendix for some tips on some useful tools installed on the VM.

### Part 1: classic stack overflow

The source and compiled binary for the program are in `/home/user/Exercise-1.1`. Take a look at the compiled binary: it is *setuid root*, so it will have super-user permissions when it executes and can do anything on the system.

**Checkpoint 1.** How can you tell the program will run as setuid root, and how can you make a compiled program run as setuid root?

Let's examine the binary a bit more with a couple of tools.

- The `checksec.sh` script can help you see what protections are enabled (or not!) in the binary. Run this on the `noticeboard` executable.
- The RATS *Rough Auditing Tool for Security* (`rats`) can help find suspicious API calls in code. Run this on the `noticeboard.c` source.

**Checkpoint 2.** Briefly explain the output of these tools, and how the compiler flags influence the output of `checksec.sh` (look at the `Makefile` to see how `gcc` was invoked).

The program contains a buffer overflow vulnerability which you should be able to see immediately. The main task of this exercise is to demonstrate this vulnerability by exploiting the program so that it launches a shell with root privileges. Here are the suggested steps:

1. Trigger a segmentation fault in the running program.
2. Use `objdump` and the debugger GDB and to inspect the code and memory layout. For example, launch the program in GDB, and find where the data is stored on the stack, and where the return address from `main` is.
3. Gain control of the instruction pointer. Do that by demonstrating how to overwrite the return address of `main` with any address of your choosing, e.g., `0xBADCODE`. Trace the execution of the program with the attack argument to check that this address is jumped to.
4. Direct execution towards some *shellcode* to execute a root shell. The shellcode could be passed in the overflow argument to execute on the stack, or set in some other way (e.g., an environment variable).

For step 4, you will need to find or write some shellcode. Samples can easily be found [online](#) but we encourage you to have a go at writing your own. A partial example was given in the lectures.

**Checkpoint 3.** Explain what your shellcode does and how you made it.

**Checkpoint 4.** Explain how your exploit works.

Finally, you should give a patch which fixes the vulnerability in the source program. Take a copy of the program, edit it and re-compile, then verify that it indeed prevents your exploit.

**Checkpoint 5.** Provide your patch to fix the notice board program (use `diff -c <oldprogram> <newprogram>`).

## Part 2: another vulnerability

Rather than patching the vulnerability, the notice board program can be compiled with some of the standard overflow protection mechanisms provided by the compiler.

Consider next the new version of the program given in `/home/user/Exercise-1.2` on the VM. The programmer of this version read on a programming blog that global variables are dangerous, so the code has been re-written to avoid using them.

Unfortunately, this new version still has vulnerabilities, especially if it runs as `setuid root`. Take a look again at the binary and source code.

**Checkpoint 1.** Identify the security flaw in the new version of `noticeboard.c`; explain what it allows and demonstrate an exploit that compromises the standard system security.

In fact, the code even be exploited to give a root shell. There is more than one way to do it, but it is a bit intricate. If you want to give it a go, feel free to try (we'll try and help!)

**Checkpoint 2 (optional).** Briefly, explain how your root shell exploit works.

**Checkpoint 3.** Give a patch which fixes the second version of `noticeboard.c`.

## Exercise 2

To avoid low-level vulnerabilities that lurk when C programming, the notice board program has now been re-written in Java. It takes the form of a *server* (which writes the messages) and a *client* program (which sends notices to the board). The client and server use a very simple message format to communicate.

Unfortunately, there is a vulnerability in this Java version of the program. A malicious user can cause the server program to crash. Since the program works over the network, this attack can be conducted remotely.

You should find this vulnerability, create an exploit to demonstrate it, patch the server program and then verify your patch prevents your exploit.

**Checkpoint 1.** Explain the format of the messages sent by the client.

**Checkpoint 2.** Provide a program (or shell script) which crashes the server remotely.

**Checkpoint 3.** Give a patch to fix the problem(s).

## Exercise 3 (Optional)

In this exercise we look at a real flaw in an old version of the [OpenSSL](#) library. OpenSSL is the main open source implementation of SSL/TLS, and is a ubiquitous part of many applications that use secure communication on the Internet.

The flaw could have been noticed by attempted attacks, in this scenario:

- A company's web servers seem to be crashing a lot.
- The logs indicate that something is causing memory to be corrupted; it seems like it happens whenever certain X.509 certificates are presented to the server.

We have provided the source code for the relevant version of OpenSSL on the virtual machine in the directory `/home/user/Exercise-3`, for you to look at. The vulnerability is inside the C function `asn1_d2i_read_bio`.

**Checkpoint 1.** Identify the security flaw in the code, and provide the relevant CVE number.

**Checkpoint 2.** Briefly summarise the problem and explain and why it is a security flaw.

**Checkpoint 3.** Give a recommendation for a way to repair the problem.

**Checkpoint 4 (very optional).** Build a *proof-of-concept* to demonstrate the security flaw and explain how it might be exploited; check that your repair (or the current released version) prevents your attack.

## Appendix: Useful commands

### General linux

**make** Builds your code

**gcc** <source-code> C compiler

-o <program> to name the executable, else it will be called **a.out**

-O 0 disables optimisations, giving simpler code

**objdump** <program> Lets you view information about a program

-x to see the headers

-d to disassemble

**strace** <program> Runs the program and shows systemcalls as they are made  
(usefull for debugging shellcode)

**nc** <host> <port> The netcat program. Useful for sending and listening to data going between ports. By default writes.

-l to listen

**ip addr** See what ip address the virtual machine has

### GDB

GDB is the go-to command line debugger. For some primitive GUI support, you can try running it inside Emacs (**M-x gdb**) or with a curses UI, **gdb --tui**.

**gdb** <program> Runs GDB on your program

-x <gdbinit> lets you run a script when GDB is first loaded

Inside GDB, the following commmands are useful:

**run** [ <args> ] [ < <input> ] Run your program with optional args and input

Can use backticks in the args to run an external command (such as **perl -e 'print "A"x9001'**)

**set args** [ <args> ] [ < <input> ] Specifies the args for the **run** command automatically

**awatch** <address> Sticks a read/write watch point whenever the memory at the address is accessed

**b <break-point>** Set a break point at a memory location (which can be a function name, e.g., **b main** or a de-referenced pointer, as **b \*0x123456**)

**c** Continue

**si** Step instruction

**x/32x \$ebp** Prints the memory as hexadecimal ints for the 32 bytes at the memory pointed to by the register **\$ebp**

**x/8i \$eip** Prints the memory as disassembled instructions the 8 bytes at the memory pointed to by the register **\$eip**

**p/x \$ebp** Print the value in the register **\$ebp**

**disas** Shows the disassembly for wherever the instruction pointer is

**list** Shows where you are in the source code (if debugging data is on)

**help** The manual!

## Radare

Radare is a *really* powerful dissembler framework but the documentation is non-existent. It is worth learning but not required for these labs. Here are a couple of commands which help in constructing shellcode.

**rasm2 "nop;nop;nop"** Gives you the bytecode for three **nop** instructions

**rasm2 -f <file> -a x86.as -C** Gives you the bytecode for assembler instructions in **<file>**, avoiding zeros and printing the result as a C-formatted string

## Submission instructions

Download a copy of the text file [checkpoints.md](http://www.inf.ed.ac.uk/teaching/courses/sp/2013/labs/lab1/checkpoints.md) from

<http://www.inf.ed.ac.uk/teaching/courses/sp/2013/labs/lab1/checkpoints.md>

and edit it to insert your answers. Submit it with the command:

```
submit sp lab1 checkpoints.md
```

We will give some feedback about the submissions at the next lab session.

*Joseph Hallett and David Aspinall, January 2014.*